

Lab 1

The target machine : RISC-V

Objective

- Be familiar with the RISC-V instruction set.
- Install the RISC-V toolchain and simulator.
- Understand how it executes on the RISC-V processor with the help of a simulator.
- Write simple programs, assemble, execute.

1.1 The RISC-V processor, instruction set, simulator

EXERCISE #1 ► Lab preparation

If you haven't done so already, follow the instructions to install `riscv-xxx-gcc` and `spike` on your machine (see `INSTALL.md` file).

EXERCISE #2 ► RISC-V C-compiler and simulator, first test

In the directory `TP01/riscv/`:

- Compile the provided file `ex1.c` with:
`riscv64-unknown-elf-gcc ex1.c -o ex1.riscv`
It produces a RISC-V binary named `ex1.riscv`.
- Execute the binary with the RISC-V simulator:
`spike pk ex1.riscv`
This should print:
`bb1 loader`
`42`
If you get a runtime exception, try running `spike -m100 pk ex1.riscv` instead: this limits the RAM usage of spike to 100 MB (the default is 2 GB).
- The corresponding RISC-V code can be obtained in a more readable format by:
`riscv64-unknown-elf-gcc ex1.c -S -o ex1.s -fverbose-asm`
(have a look at the generated `.s` file!)

The objective of this sequence of labs is to design **our own (subset of) C compiler for RISC-V**.

EXERCISE #3 ► Documents

Some documentation can be found in the RISC-V ISA (`riscv_isa.pdf` on the course webpage).

<https://compil-lyon.gitlabpages.inria.fr/>

In the architecture course, you already saw a version of the target machine RISC-V. The instruction set is depicted in Appendix ??.

1.1.1 Hand exercises

EXERCISE #4 ► TD

On paper, write (in RISC-V assembly language) a program which initializes the t_0 register to 1 and increments it until it becomes equal to 8.

EXERCISE #5 ► TD : sum

Write a program in RISC-V assembly that computes the sum of the 10 first positive integers (excluded 10).

1.1.2 Assembling, disassembling

EXERCISE #6 ► Hand assembling, simulation of the hex code

Assemble by hand (on paper) the instructions:

```

1      .globl main
2 main:
3      addi a0, a0, 1
4      bne a0, a0, main
5 end:
6      ret

```

You will need the set of instructions of the RISC-V machine and their associated opcode. All the info is in the ISA documentation.

To check your solution (**after** you did the job manually), you can redo the assembly using the toolchain:

```
riscv64-unknown-elf-as -march=rv64g asshand.s -o asshand.o
```

asshand.o is an ELF file which contains both the compiled code and some metadata (you can try `hexdump asshand.o` to view its content, but it's rather large and unreadable). The tool `objdump` allows extracting the code section from the executable, and show the binary code next to its disassembled version:

```
riscv64-unknown-elf-objdump -d asshand.o
```

Check that the output is consistent with what you found manually.

EXERCISE #7 ► Hand disassembling

Guess a RISC-V program that assembles itself into :

Listing 1.1: disass.lst

```
disass.o: format de fichier elf64-littleriscv
```

Déassemblage de la section .text:

```

0000000000000000 <main>:
0: 00128313  xx
4: ffdff06f  yy
8: 00008067  zz

```

From now on, we are going to write programs using an easier approach. We are going to write instructions using the RISC-V assembly.

1.2 RISC-V Simulator

EXERCISE #8 ► Execution and debugging

See <https://www.lowrisc.org/docs/tagged-memory-v0.1/spike/> for details on the Spike simulator.

`test_print.s` is a small but complete example using Risc-V assembly. It uses the `println_string`, `print_int`, `print_char` and `newline` functions provided to you in `libprint.s`. Each function can be called with `call print_...` and prints the content of register `a0` (`call newline` takes no input and prints a new-line character).

1. First test assembling and simulation on the file `test_print.s`:

```
riscv64-unknown-elf-as -march=rv64g test_print.s -o test_print.o
```
2. Optionally, run `riscv64-unknown-elf-objdump -D` as in previous exercise. The `-D` option shows all sections, including `.rodata`.
3. The `libprint.s` library must be assembled too:

- ```
riscv64-unknown-elf-as -march=rv64g libprint.s -o libprint.o
```
- We now link these files together to get an executable:  

```
riscv64-unknown-elf-gcc test_print.o libprint.o -o test_print
```

The generated `test_print` file should be executable, but since it uses the Risc-V ISA, we can't execute it natively (try `./test_print`, you'll get an error like `Exec format error`).
  - Run the simulator:  

```
spike pk ./test_print
```

The output should look like:  

```
bbl loader
HI MIF08!
42
a
```

The first line comes from the simulator itself, the next two come from the `println_string`, `print_int` and `print_char` calls in the assembly code.
  - We can also view the instructions while they are executed:  

```
spike -l pk ./test_print
```

Unfortunately, this shows all the instructions in `pk` (Proxy Kernel, a kind of mini operating system), and is mostly unusable. Alternatively, we can run a step-by-step simulation starting from a given symbol. To run the instructions in `main`, we first get the address of `main` in the executable:  

```
$ riscv64-unknown-elf-nm test_print | grep main
0000000000001014c T main
```

This means: `main` is a symbol defined in the `.text` section (T in the middle column), it is global (capital T), and its address is 1014c (you may not have the same address). Now, run `spike` in debug mode (`-d`) and execute code up to this address (until `pc 0 1014c`, i.e. "Until the program counter of core 0 reaches 1014c"). Press Return to move to the next instruction and `q` to quit:  

```
$ spike -d pk ./test_print
: until pc 0 1014c
bbl loader
:
core 0: 0x0000000000001014c (0xff010113) addi sp, sp, -16
:
core 0: 0x00000000000010150 (0x00113423) sd ra, 8(sp)
:
core 0: 0x00000000000010154 (0x0000e517) auipc a0, 0xe
:
core 0: 0x00000000000010158 (0x41450513) addi a0, a0, 1044
: q
$
```

**Remark:** For your labs, you may want to assemble and link with a single command (which can also do the compilation if you provide `.c` files on the command-line):

```
riscv64-unknown-elf-gcc -march=rv64g libprint.s test_print.s -o main
```

In real-life, people run compilation+assembly and link as two different commands, but use a build system like a `Makefile` to re-run only the right commands.

### EXERCISE #9 ► Algo in RISC-V assembly

Write (in `minmax.s`) a program in RISC-V assembly that computes the min of two integers, and stores the result in a precise location of the memory that has the label `min`. Try with different values. We use 64 bits of memory to store ints, i.e. use `.dword` directive and `ld` and `sd` instructions.

### EXERCISE #10 ► (Advanced) Algo in RISC-V assembly

Write and execute the following programs in assembly:

- Count the number of non-nul bits of a given integer, print the result.

- Draw squares and triangles of stars (character '\*') of size  $n$ ,  $n$  being stored somewhere in memory.

Examples:

$n=3$  square:

\*\*\*

\*\*\*

\*\*\*

$n=3$  triangle:

  \*

 \* \*

\* \* \*

The function `print_char` expects the ASCII encoding of the character you want to print.